

```
true
```

```
• begin
•   using ExtendableSparse
•   using SparseArrays
•   using ForwardDiff
•   using SparseDiffTools
•   using SparsityDetection
•   using Symbolics
•   using VoronoiFVM
•   using PlutoUI
•   using PyPlot
•   PyPlot.svg(true);
• end
```

Table of Contents

Ways to calculate sparse Jacobians

The methods

Straightforward Jacobian calculation

Symbolics.jl for sparsity detection

SparsityDetection.jl for sparsity detection

Assembly from local Jacobians

Assembly from local:VoronoiFVM.jl

Scaling comparison

Discussion

```
• TableOfContents()
```

Ways to calculate sparse Jacobians

We investigate different ways to assemble sparse Jacobians for a finite difference operator $A_h(u)$ discretizing the differential operator

$$A(u) = -\Delta u^2 + u^2 - 1$$

in $\Omega = (0, 1)$ with homogeneous Neumann boundary conditions on an equidistant grid of n points. We only discuss approaches which can be generalized to higher space dimensions and unstructured grids, so e.g. banded and tridiagonal matrix structures are not discussed.

This is the nonlinear finite difference operator:

```

• function A_h!(y,u)
•     n=length(u)
•     h=1/(n-1)
•     y[1]=(u[1]^2-1)*0.5*h
•     for i=2:n-1
•         y[i]=(u[i]^2-1.0)*h
•     end
•     y[end]=(u[end]^2-1)*0.5*h
•
•     for i=1:n-1
•         du=(u[i+1]^2-u[i]^2)/h
•         y[i+1]+=du
•         y[i]-=du
•     end
• end;

```

Number of discretization points for testing:

```
• n_test=5;
```

The methods

Straightforward Jacobian calculation

Just take the operator and calculate the Jacobian as a sparse matrix. Here we use the `ExtendableSparse.jl` package in order to avoid cumbersome handling of intermediate coordinate format data.

As we will see in the timing test, the complexity of this operation is $O(n^2)$

```

• function straightforward(n)
•     X=ones(n)
•     Y=ones(n)
•     jac = ExtendableSparseMatrix(n, n)
•     t=@elapsed begin
•         ForwardDiff.jacobian!(jac, A_h!, X, Y)
•         flush!(jac)
•     end
•     jac.cscmatrix,t
• end;

```

```

(5×5 SparseMatrixCSC{Float64, Int64} with 13 stored entries:, 0.992594)
 8.25  -8.0   .   .   .
-8.0   16.5  -8.0  .   .
 .   -8.0   16.5  -8.0  .
 .   .   -8.0   16.5  -8.0
 .   .   .   -8.0   8.25

```

```
• straightforward(n_test)
```

Symbolics.jl for sparsity detection

While in the particular case we the sparsity pattern of the operator – it is tridiagonal – we assume it is not known and we use Symbolics.jl to detect the sparsity pattern and convey the sparsity information to ForwardDiff.

For this purpose we use the multiple dispatch feature of Julia and call `A_h!` with symbolic data, resulting in a symbolic representation of the operator which can be investigated for sparsity in $O(n)$ time.

```
• function Symbolics.jacobian_sparsity(op!,output::Array{T},input::Array{T})
  where T<:Number
•   eqs=similar(output,Num)
•   vars=[Symbolics.variable(i) for i in eachindex(input)]
•   op!(eqs,vars)
•   Symbolics.jacobian_sparsity(eqs,vars)
• end
```

```
• function symbolics(n)
•   input = ones(n)
•   output = similar(input)
•   tdetect=@elapsed begin
•       sparsity_pattern=Symbolics.jacobian_sparsity(A_h!,output,input)
•       jac=Float64.(sparsity_pattern)
•       colors = matrix_colors(jac)
•   end
•   tassemble=@elapsed begin
•       forwarddiff_color_jacobian!(jac, A_h!, input, colorvec = colors)
•   end
•   jac,tdetect,tassemble
• end;
```

```
(5×5 SparseMatrixCSC{Float64, Int64} with 13 stored entries:, 0.000492987, 4.0342e-!
 8.25  -8.0   .   .   .
-8.0  16.5  -8.0   .   .
.   -8.0  16.5  -8.0   .
.   .   -8.0  16.5  -8.0
.   .   .   -8.0  8.25
```

```
• symbolics(n_test)
```

SparsityDetection.jl for sparsity detection

Before the advent of Symbolics.jl, SparsityDetection.jl (now deprecated) provided a method for sparsity detection of nonlinear operators in $O(n)$ time which we test here for comprehensiveness.

```

• function sparsitydetect(n)
•     input = rand(n)
•     output = similar(input)
•     tdetect=@elapsed begin
•         sparsity_pattern = jacobian_sparsity(A_h!,output,input)
•         jac = Float64.(sparsity_pattern)
•         colors = matrix_colors(jac)
•     end
•     u=ones(n)
•     tassemble=@elapsed begin
•         forwarddiff_color_jacobian!(jac, A_h!, u, colorvec = colors)
•     end
•     jac,tdetect,tassemble
• end;

```

```

(5×5 SparseMatrixCSC{Float64, Int64} with 13 stored entries:, 0.000837247, 4.2508e-!
 8.25  -8.0   .   .   .
-8.0   16.5  -8.0   .   .
.   -8.0   16.5  -8.0   .
.   .   -8.0   16.5  -8.0
.   .   .   -8.0   8.25

```

```

• sparsitydetect(n_test)

```

Assembly from local Jacobians

Usually, for finite difference and finite volume methods, the sparsity structure is defined by the grid topology, and the operator is essentially a superposition of local contributions translated in space. So we can assemble the global Jacobi matrix from local contributions, very much like in classical finite element assembly. The local Jacobi matrix can be obtained without the need to detect its sparsity (though for larger coupled systems of PDEs this possibility appears to be worth to be investigated)

For this purpose, we define the local contributions as mutating Julia functions, ready to be generalized for systems of PDEs.

```

• flux!(f,u) = f[1]=u[1]^2-u[2]^2;

```

```

• reaction!(f,u) = f[1]=u[1]^2-1.0;

```

These can be passed to a function which calculates the operator application and assembles the Jacobi matrix at once as it is needed in Newton's method. It is important to ensure that there are no allocations in the inner loop.

```

• function A_Jac_h!(y,m,u,flux,reaction)
•     n=length(u)
•     h=1/(n-1)
•
•     Y=zeros(1)
•     UK=zeros(1)
•     UKL=zeros(2)
•
•     # Provide space for results
•     result_rea = DiffResults.DiffResult(zeros(1), zeros(1,1))
•     result_flx  = DiffResults.DiffResult(zeros(1), zeros(1,2))
•
•     # Fix the Jacobian configuration for ForwardDiff
•     cfg_rea     = ForwardDiff.JacobianConfig(reaction, Y, UK,
ForwardDiff.Chunk(UK,1))
•     cfg_flx     = ForwardDiff.JacobianConfig(flux, Y,
UKL,ForwardDiff.Chunk(UKL,1))
•
•
•     for i=1:n
•         fac=h
•         if i==1 || i==n
•             fac=0.5*h
•         end
•         UK[1]=u[i]
•         # we could use ForwardDiff.jacobian! here, but that is not allocation
•     free
•
•         ForwardDiff.vector_mode_jacobian!(result_rea,reaction,Y,UK,cfg_rea)
•         y[i]=DiffResults.value(result_rea)[1]*fac
•         m[i,i] += DiffResults.jacobian(result_rea)[1,1]*fac
•     end
•
•     for i=1:n-1
•         UKL[1]=u[i]
•         UKL[2]=u[i+1]
•         ForwardDiff.vector_mode_jacobian!(result_flx,flux,Y,UKL,cfg_flx)
•         du=DiffResults.value(result_flx)
•         ddu=DiffResults.jacobian(result_flx)
•         y[i]+=du[1]/h
•         y[i+1]-=du[1]/h
•         m[i,i]+=ddu[1,1]/h
•         m[i,i+1]+=ddu[1,2]/h
•         m[i+1,i]-=ddu[1,1]/h
•         m[i+1,i+1]-=ddu[1,2]/h
•     end
• end;

```

For the allocations in `ForwardDiff.jacobian!`, see [ForwardDiff.jl#516](#).

```

• function assemblefromlocal(n)
•     u=ones(n)
•     jac = ExtendableSparseMatrix(n,n);
•     y=ones(n)
•     t=@elapsed begin
•         A_Jac_h!(y,jac,u,flux!,reaction!)
•         flush!(jac)
•     end
•     jac.cscmatrix,t
• end;

```

```
(5×5 SparseMatrixCSC{Float64, Int64} with 13 stored entries:, 0.543577)
 8.25  -8.0   .   .   .
-8.0  16.5 -8.0   .   .
 .   -8.0  16.5 -8.0   .
 .   .   -8.0  16.5 -8.0
 .   .   .   -8.0  8.25
```

- `assemblefromlocal(n_test)`

Assembly from local: VoronoiFVM.jl

The previous approach is at the core of VoronoiFVM.jl, which works for PDE systems on unstructured grids in 1/2/3D.

`vfvm_flux!` (generic function with 1 method)

- `vfvm_flux!(f,u,edge)=f[1]=u[1,1]^2-u[1,2]^2`

`vfvm_reaction!` (generic function with 1 method)

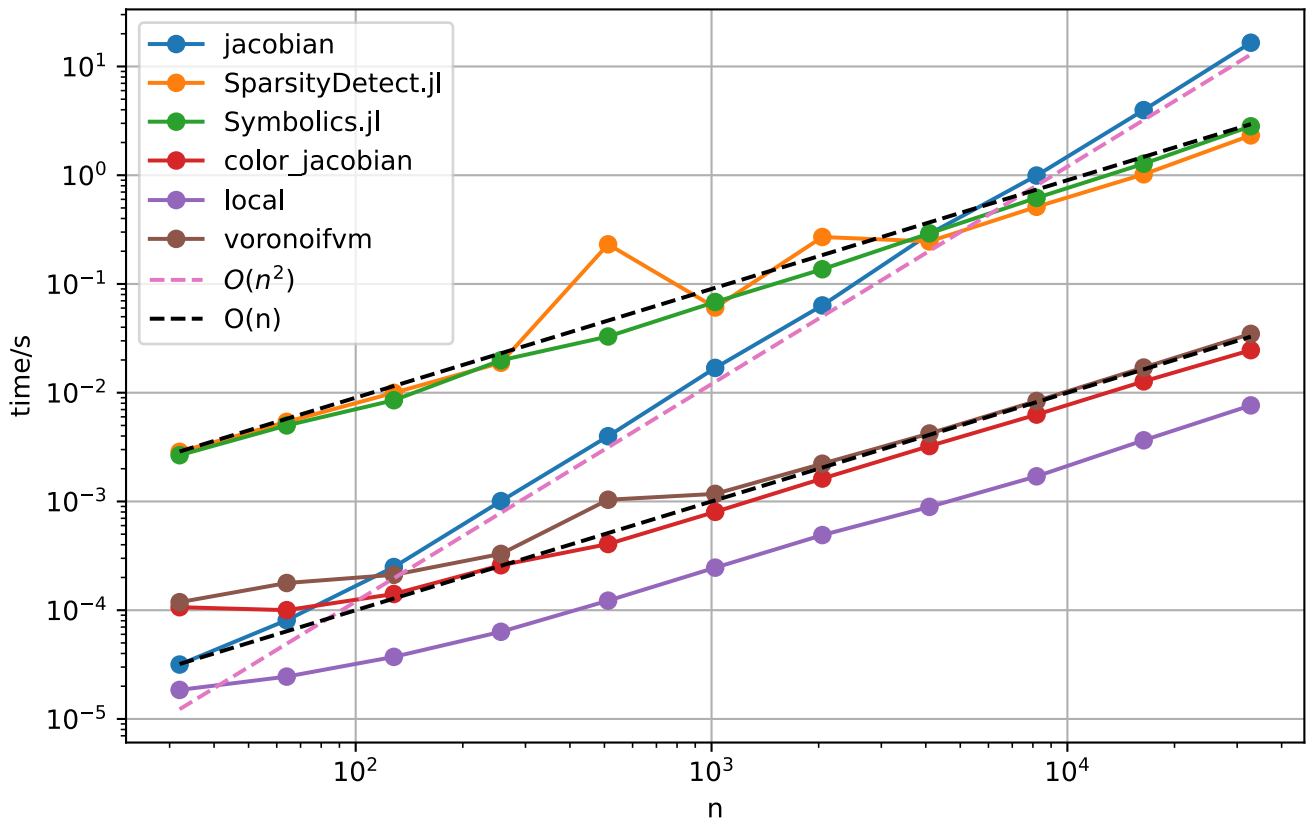
- `vfvm_reaction!(f,u,node)=f[1]=u[1]^2-1.0`

```
• function voronoifvm(n)
•     h=1.0/convert(Float64,n-1)
•     X=collect(0:h:1)
•     grid=VoronoiFVM.Grid(X)
•     physics=VoronoiFVM.Physics(flux=vfvm_flux!,reaction=vfvm_reaction!)
•     sys=VoronoiFVM.System(grid,physics,unknown_storage=:dense)
•     enable_species!(sys,1,[1])
•     solution=unknowns(sys,inival=1.0)
•     oldsol=unknowns(sys,inival=1.0)
•     residual=unknowns(sys,inival=1.0)
•     timestep=1.0e100
•
•     t=@elapsed begin
•         VoronoiFVM.eval_and_assemble(sys,solution,oldsol,residual,timestep)
•         flush!(sys.matrix)
•     end
•     sys.matrix.cscmatrix,t
• end;
```

```
(5×5 SparseMatrixCSC{Float64, Int64} with 13 stored entries:, 3.10021)
 8.25  -8.0   .   .   .
-8.0  16.5 -8.0   .   .
 .   -8.0  16.5 -8.0   .
 .   .   -8.0  16.5 -8.0
 .   .   .   -8.0  8.25
```

- `voronoifvm(5)`

Scaling comparison



```
• plot_scaling(powmax=15)
```

Discussion

- As expected, the straightforward method shows $O(n^2)$ complexity.
- Julia's sparsity detection has $O(n)$ complexity but is around two orders of magnitudes slower than `color_jacobian` used to calculate the actual Jacobian based on the sparsity detected.
- Assembly from local Jacobians as implemented here is $O(n)$ and the fastest option. It does not need the sparsity detection step.
- Assembly from local Jacobians in `VoronoiFVM.jl` is $O(n)$ but performs worse than the sample implementation provided and slightly worse than `color_jacobian`. This is due to the bookkeeping for the discretization grid and other overheads which come from the generic nature of that package. It does not need the sparsity detection step as well.

